

Distributed Data Analytics

Lab 7

Submitted By- Abdullah Al Foysal

Id:277458

Exercise 1: Convolutional Neural Networks with Relu activation function

```
import pickle
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn import preprocessing
#import numpy as np
import random
#import tensorflow as tf
from sklearn.preprocessing import LabelBinarizer

# This function returns the number of labels in the cifar10 dataset

def _load_label_names():

    return ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# This function takes the batches from the cifar10 folder and convert each batch size 3072 to the transformation of depth,length,
def load_cifar10_batch(cifar10_dataset_folder_path, batch_id):

    with open(cifar10_dataset_folder_path + '/data_batch_' + str(batch_id), mode='rb') as file:
        batch = pickle.load(file, encoding='latin1')

    features = batch['data'].reshape((len(batch['data']), 3, 32, 32)).transpose(0, 2, 3, 1)
    labels = batch['labels']

    return features, labels

#This function is used to display the image statistics based on the batch_id and sample_id
#If the batch_id and sample_id is 4 and 1 respectively, it prints the first mage from training batch 4 along with its statistics
```

```

def display_stats(cifar10_dataset_folder_path, batch_id, sample_id):

    batch_ids = list(range(1, 6))

    if batch_id not in batch_ids:
        print('Batch Id out of Range. Possible Batch Ids: {}'.format(batch_ids))
        return None

    features, labels = load_cifar10_batch(cifar10_dataset_folder_path, batch_id)

    if not (0 <= sample_id < len(features)):
        print('{} samples in batch {}. {} is out of range.'.format(len(features), batch_id, sample_id))
        return None

    print('\nStats of batch {}'.format(batch_id))
    print('Samples: {}'.format(len(features)))
    print('Label Counts: {}'.format(dict(zip(*np.unique(labels, return_counts=True))))))
    print('First 20 Labels: {}'.format(labels[:20]))

    sample_image = features[sample_id]
    sample_label = labels[sample_id]
    label_names = _load_label_names()

    print('\nExample of Image {}'.format(sample_id))
    print('Image - Min Value: {} Max Value: {}'.format(sample_image.min(), sample_image.max()))
    print('Image - Shape: {}'.format(sample_image.shape))
    print('Label - Label Id: {} Name: {}'.format(sample_label, label_names[sample_label]))
    plt.axis('off')
    plt.imshow(sample_image)

#This function is used to do preprocessing stuff like normalizing the features and one hot encoding the labels and the result is
def _preprocess_and_save(normalize, one_hot_encode, features, labels, filename):

    features = normalize(features)
    labels = one_hot_encode(labels)

    pickle.dump((features, labels), open(filename, 'wb'))

```

```

def preprocess_and_save_data(cifar10_dataset_folder_path, normalize, one_hot_encode):

    n_batches = 5
    valid_features = []
    valid_labels = []

    for batch_i in range(1, n_batches + 1):
        features, labels = load_cifar10_batch(cifar10_dataset_folder_path, batch_i)
        validation_count = int(len(features) * 0.1)

        # Preprocess and save a batch of training data
        _preprocess_and_save(
            normalize,
            one_hot_encode,
            features[:-validation_count],
            labels[:-validation_count],
            'preprocess_batch_' + str(batch_i) + '.p')

        # Use a portion of training batch for validation
        valid_features.extend(features[-validation_count:])
        valid_labels.extend(labels[-validation_count:])

    # Preprocess and Save all validation data
    _preprocess_and_save(
        normalize,
        one_hot_encode,
        np.array(valid_features),
        np.array(valid_labels),
        'preprocess_validation.p')

    with open(cifar10_dataset_folder_path + '/test_batch', mode='rb') as file:
        batch = pickle.load(file, encoding='latin1')

```

```

# Load the testing data

test_features = batch['data'].reshape((len(batch['data']), 3, 32, 32)).transpose(0, 2, 3, 1)
test_labels = batch['labels']

# Preprocess and Save all testing data
_preprocess_and_save(
    normalize,
    one_hot_encode,
    np.array(test_features),
    np.array(test_labels),
    'preprocess_test.p')

def batch_features_labels(features, labels, batch_size):

    for start in range(0, len(features), batch_size):
        end = min(start + batch_size, len(features))
        yield features[start:end], labels[start:end]

#Load the preprocessed training batch
def load_preprocess_training_batch(batch_id, batch_size):

    filename = 'preprocess_batch_' + str(batch_id) + '.p'
    features, labels = pickle.load(open(filename, mode='rb'))

    # Return the training data in batches of size <batch_size> or less
    return batch_features_labels(features, labels, batch_size)

#Display the prediction for the random test data based on the model saved after training it.
def display_image_predictions(features, labels, predictions):
    #it is the number of classes
    n_classes = 10
    label_names = _load_label_names()
    label_binarizer = LabelBinarizer()
    label_binarizer.fit(range(n_classes))
    label_ids = label_binarizer.inverse_transform(np.array(labels))

    fig, axes = plt.subplots(nrows=4, ncols=2)
    fig.tight_layout()
    fig.suptitle('Softmax Predictions', fontsize=20, y=1.1)
    #Top 3 probability predictions are shown for each random predicted test data.
    n_predictions = 3
    margin = 0.05
    ind = np.arange(n_predictions)
    width = (1. - 2. * margin) / n_predictions

    for image_i, (feature, label_id, pred_indicies, pred_values) in enumerate(zip(features, label_ids, predictions.indices, predi
        pred_names = [label_names[pred_i] for pred_i in pred_indicies]
        correct_name = label_names[label_id]

        axes[image_i][0].imshow(feature)
        axes[image_i][0].set_title(correct_name)
        axes[image_i][0].set_axis_off()

        axes[image_i][1].barh(ind + margin, pred_values[::-1], width)
        axes[image_i][1].set_yticks(ind + margin)
        axes[image_i][1].set_yticklabels(pred_names[::-1])
        axes[image_i][1].set_xticks([0, 0.5, 1.0])

#This is the folder path where i downloaded and extracted my cifar10 dataset
cifar10_dataset_folder_path = 'D:\\\\Uni Hildesheim\\\\4th semester\\\\DDA Lab\\\\lab 7\\\\cifar-10-python\\\\cifar-10-batches-py'
def normalize(x):

    maximum = np.max(x)
    minimum = np.min(x)
    return (x - minimum) / (maximum - minimum)

```



```

#Display the prediction for the random test data based on the model saved after training it.
def display_image_predictions(features, labels, predictions):
    #it is the number of classes
    n_classes = 10
    label_names = _load_label_names()
    label_binarizer = LabelBinarizer()
    label_binarizer.fit(range(n_classes))
    label_ids = label_binarizer.inverse_transform(np.array(labels))

    fig, axes = plt.subplots(nrows=4, ncols=2)
    fig.tight_layout()
    fig.suptitle('Softmax Predictions', fontsize=20, y=1.1)
    #Top 3 probability predictions are shown for each random predicted test data.
    n_predictions = 3
    margin = 0.05
    ind = np.arange(n_predictions)
    width = (1. - 2. * margin) / n_predictions

    for image_i, (feature, label_id, pred_indicies, pred_values) in enumerate(zip(features, label_ids, predictions.indices, predi
        pred_names = [label_names[pred_i] for pred_i in pred_indicies]
        correct_name = label_names[label_id]

        axes[image_i][0].imshow(feature)
        axes[image_i][0].set_title(correct_name)
        axes[image_i][0].set_axis_off()

        axes[image_i][1].barh(ind + margin, pred_values[::-1], width)
        axes[image_i][1].set_yticks(ind + margin)
        axes[image_i][1].set_yticklabels(pred_names[::-1])
        axes[image_i][1].set_xticks([0, 0.5, 1.0])

#This is the folder path where i downloaded and extracted my cifar10 dataset
cifar10_dataset_folder_path = 'D:\\Uni Hildesheim\\4th semester\\DDA Lab\\lab 7\\cifar-10-python\\cifar-10-batches-py'
def normalize(x):

    maximum = np.max(x)
    minimum = np.min(x)
    return (x - minimum) / (maximum - minimum)

```

```

#Defining the convolutional neural network with number of output, kernel size, strides, and pool side and padding
def conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize, pool_strides):

```

```

    input_depth = x_tensor.get_shape().as_list()[-1]
    #weights
    W = tf.Variable(tf.random_normal(
        [conv_ksize[0], conv_ksize[1], input_depth, conv_num_outputs],
        stddev=0.1
    ))
    #bias
    b = tf.Variable(tf.zeros(conv_num_outputs))
    conv = tf.nn.conv2d(x_tensor, W, [1, conv_strides[0], conv_strides[1], 1], 'SAME') + b
    conv = tf.nn.relu(conv)
    return tf.nn.max_pool(
        conv,
        [1, pool_ksize[0], pool_ksize[1], 1],
        [1, pool_strides[0], pool_strides[1], 1],
        'SAME')

```

```

def flatten(x_tensor):

```

```

    shape = x_tensor.get_shape().as_list()
    return tf.reshape(x_tensor, [-1, np.prod(shape[1:])])

```

```

#Defined a Fully connected layer with relu activation function

```

```

def fully_conn(x_tensor, num_outputs):

```

```

    shape = x_tensor.get_shape().as_list()
    W = tf.Variable(tf.random_normal([shape[-1], num_outputs], stddev=0.1))
    b = tf.Variable(tf.zeros(num_outputs)) + 0.11
    return tf.nn.relu(tf.add(tf.matmul(x_tensor, W), b))

```

```

#defining the function for output of each network

```

```

def output(x_tensor, num_outputs):

```

```

    shape = x_tensor.get_shape().as_list()
    W = tf.Variable(tf.random_normal([shape[-1], num_outputs]))
    b = tf.Variable(tf.zeros(num_outputs))
    return tf.add(tf.matmul(x_tensor, W), b)

```

```

def conv_net(x):

    # conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize, pool_strides)
    tmp = conv2d_maxpool(x, 64, [3, 3], [1, 1], [3, 3], [2, 2])

    tmp = flatten(tmp)
    #one fully connected layer fc1
    tmp = fully_conn(tmp, 384)

    #second fully connected layer fc2
    tmp = fully_conn(tmp, 192)

    # TODO: return output
    return output(tmp, 10)

tf.reset_default_graph()

# Inputs
x = neural_net_image_input((32, 32, 3))
y = neural_net_label_input(10)

# Model
logits = conv_net(x)

# Name logits Tensor, so that it can be loaded from disk after training
logits = tf.identity(logits, name='logits')

# Loss and Optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))
optimizer = tf.train.AdamOptimizer().minimize(cost)

# Accuracy
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')
tf.summary.histogram("loss", cost)
tf.summary.scalar("loss", cost)
# Create a summary to monitor accuracy tensor
tf.summary.histogram("accuracy", accuracy)
tf.summary.scalar("accuracy", accuracy)
# Merge all summaries into a single op
merged_summary_op = tf.summary.merge_all()

def train_neural_network(session, optimizer, cost, merged_summary_op, feature_batch, label_batch, epoch):

    train_opt, train_cost, summary = session.run([optimizer, cost, merged_summary_op], feed_dict={x: feature_batch, y: label_batch})
    summary_writer.add_summary(summary, epoch)

def print_stats(session, feature_batch, label_batch, cost, accuracy):

    global valid_features, valid_labels
    validation_accuracy = session.run(
        accuracy,
        feed_dict={
            x: valid_features,
            y: valid_labels,
        }
    )
    cost = session.run(
        cost,
        feed_dict={
            x: feature_batch,
            y: label_batch,
        }
    )
    print('Cost = {0} - Validation Accuracy = {1}'.format(cost, validation_accuracy))

# Parameters
epochs = 5
batch_size = 1024

```



```
#This the model path where i save my trained result which can be later used for predicting the results

save_model_path = 'D:\\Uni Hildesheim\\4th semester\\DDA Lab\\lab 7\\cifar-10-python\\cifar-10-batches-py\\image_classification'

print('Training...')
with tf.Session() as sess:
    # Initializing the variables
    sess.run(tf.global_variables_initializer())
    summary_writer = tf.summary.FileWriter("D:\\Uni Hildesheim\\4th semester\\DDA Lab\\lab 7\\CNN-relu\\train\\", graph=tf.get_default_graph())

    # Training cycle
    for epoch in range(epochs):
        # Loop over all batches
        n_batches = 5
        for batch_i in range(1, n_batches + 1):
            for batch_features, batch_labels in load_preprocess_training_batch(batch_i, batch_size):
                train_neural_network(sess, optimizer, cost, merged_summary_op, batch_features, batch_labels, epoch)
            print('Epoch {:>2}, CIFAR-10 Batch {:>2}'.format(epoch + 1, batch_i), end='')
            print_stats(sess, batch_features, batch_labels, cost, accuracy)

    # Save Model
    saver = tf.train.Saver()
    save_path = saver.save(sess, save_model_path)

# # I have ran each training batch for 5 epoch, as i dont have GPU and my CPU takes more time for running. The model which i have
```

WARNING:tensorflow:From <ipython-input-1-032265a8074a>:277: softmax_cross_entropy_with_logits (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See `@(tf.nn.softmax_cross_entropy_with_logits_v2)`.

```
Training...
Epoch 1, CIFAR-10 Batch 1: Cost = 31.28889274597168 - Validation Accuracy = 0.1242000013589859
Epoch 1, CIFAR-10 Batch 2: Cost = 9.799338340759277 - Validation Accuracy = 0.1518000066280365
Epoch 1, CIFAR-10 Batch 3: Cost = 4.63309907913208 - Validation Accuracy = 0.18140000104904175
Epoch 1, CIFAR-10 Batch 4: Cost = 3.597994565963745 - Validation Accuracy = 0.21199999749660492
Epoch 1, CIFAR-10 Batch 5: Cost = 2.816688060760498 - Validation Accuracy = 0.22939999401569366
Epoch 2, CIFAR-10 Batch 1: Cost = 2.629655599594116 - Validation Accuracy = 0.28200000524520874
Epoch 2, CIFAR-10 Batch 2: Cost = 2.3413617610931396 - Validation Accuracy = 0.30079999566078186
Epoch 2, CIFAR-10 Batch 3: Cost = 2.0733025074005127 - Validation Accuracy = 0.3393999934196472
Epoch 2, CIFAR-10 Batch 4: Cost = 1.9147813320159912 - Validation Accuracy = 0.3691999912261963
Epoch 2, CIFAR-10 Batch 5: Cost = 1.7407573461532593 - Validation Accuracy = 0.3846000134944916
Epoch 3, CIFAR-10 Batch 1: Cost = 1.7929636240005493 - Validation Accuracy = 0.40459999442100525
Epoch 3, CIFAR-10 Batch 2: Cost = 1.6771470308303833 - Validation Accuracy = 0.4334000051021576
Epoch 3, CIFAR-10 Batch 3: Cost = 1.5652251243591309 - Validation Accuracy = 0.4496000111103058
Epoch 3, CIFAR-10 Batch 4: Cost = 1.4940121173858643 - Validation Accuracy = 0.46160000562667847
Epoch 3, CIFAR-10 Batch 5: Cost = 1.3686408996582031 - Validation Accuracy = 0.4729999899864197
Epoch 4, CIFAR-10 Batch 1: Cost = 1.4570103883743286 - Validation Accuracy = 0.48399999737739563
Epoch 4, CIFAR-10 Batch 2: Cost = 1.3770177364349365 - Validation Accuracy = 0.49720001220703125
Epoch 4, CIFAR-10 Batch 3: Cost = 1.3127710819244385 - Validation Accuracy = 0.5041999816894531
Epoch 4, CIFAR-10 Batch 4: Cost = 1.2679089307785034 - Validation Accuracy = 0.5175999999046326
Epoch 4, CIFAR-10 Batch 5: Cost = 1.1604394912719727 - Validation Accuracy = 0.5216000080108643
Epoch 5, CIFAR-10 Batch 1: Cost = 1.2559391260147095 - Validation Accuracy = 0.5242000222206116
Epoch 5, CIFAR-10 Batch 2: Cost = 1.2053776979446411 - Validation Accuracy = 0.5347999930381775
Epoch 5, CIFAR-10 Batch 3: Cost = 1.15341317653656 - Validation Accuracy = 0.5479999780654907
Epoch 5, CIFAR-10 Batch 4: Cost = 1.1412394046783447 - Validation Accuracy = 0.5410000085830688
Epoch 5, CIFAR-10 Batch 5: Cost = 1.038729190826416 - Validation Accuracy = 0.5550000071525574
```

I have run each training batch for 5 epoch, as I don't have GPU and my CPU takes more time for running. The model which I have saved for training will now be used back for predicting the results for any random 4 images from the test dataset. It produces random images with its

top-n (here it's 3) probability of that image belonging to those classes and testing accuracy is plotted. My average training accuracy at the end of epoch 5 is 0.55 (i.e. 55%)

```
try:
    if batch_size:
        pass
except NameError:
    batch_size = 64
#sample of images to be displayed is 4
n_samples = 4
#prediction of each image is its top 3 prediction.
top_n_predictions = 3

def test_model():

    test_features, test_labels = pickle.load(open('preprocess_test.p', mode='rb'))
    loaded_graph = tf.Graph()

    with tf.Session(graph=loaded_graph) as sess:
        # Load model
        loader = tf.train.import_meta_graph(save_model_path + '.meta')
        loader.restore(sess, save_model_path)

        # Get Tensors from Loaded model
        loaded_x = loaded_graph.get_tensor_by_name('x:0')
        loaded_y = loaded_graph.get_tensor_by_name('y:0')
        loaded_logits = loaded_graph.get_tensor_by_name('logits:0')
        loaded_acc = loaded_graph.get_tensor_by_name('accuracy:0')

        # Get accuracy in batches for memory limitations
        test_batch_acc_total = 0
        test_batch_count = 0

        for test_feature_batch, test_label_batch in batch_features_labels(test_features, test_labels, batch_size):
            test_batch_acc_total += sess.run(
                loaded_acc,
                feed_dict={loaded_x: test_feature_batch, loaded_y: test_label_batch})
            test_batch_count += 1

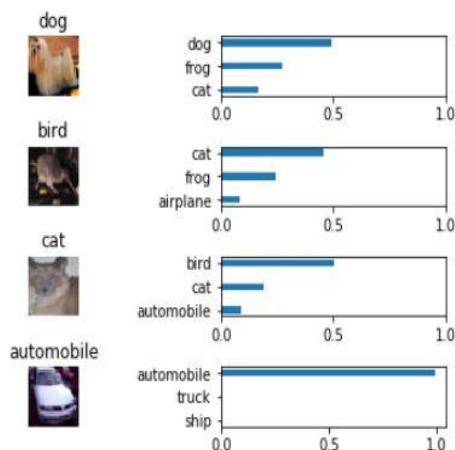
        print('Testing Accuracy: {}'.format(test_batch_acc_total/test_batch_count))

        # Print Random Samples
        random_test_features, random_test_labels = tuple(zip(*random.sample(list(zip(test_features, test_labels)), n_samples)))
        random_test_predictions = sess.run(
            tf.nn.top_k(tf.nn.softmax(loaded_logits), top_n_predictions),
            feed_dict={loaded_x: random_test_features, loaded_y: random_test_labels})
        display_image_predictions(random_test_features, random_test_labels, random_test_predictions)

test_model()
```

INFO:tensorflow:Restoring parameters from D:\Uni Hildesheim\4th semester\DDA Lab\lab 7\cifar-10-python\cifar-10-batches-py\image_classification
Testing Accuracy: 0.5525171399116516

Softmax Predictions

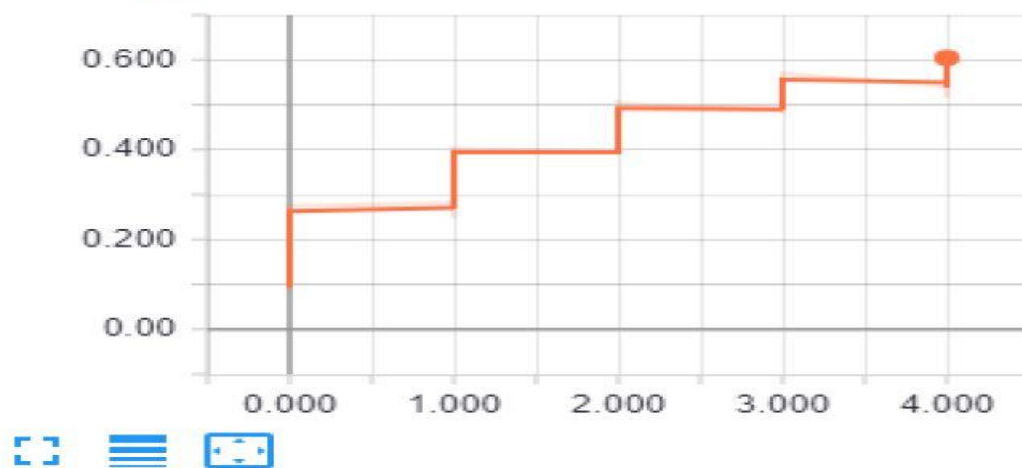


My testing accuracy was 0.552 or (55.2%) for cifar-10 dataset using elu as activation function without any batch normalization. The above are the prediction of random 4 images and top 3 prediction with respect to each images was predicted correctly or not. For this example image 1 and image 4 were predicted properly with SoftMax probability > 0.5 , for 2nd and 3rd image it was misclassified (expected bird but displayed as cat and expected cat but displayed as bird as higher probability)


Tensorboard Visualization

Scalar

accuracy_2



loss_1




Distributions

accuracy_1

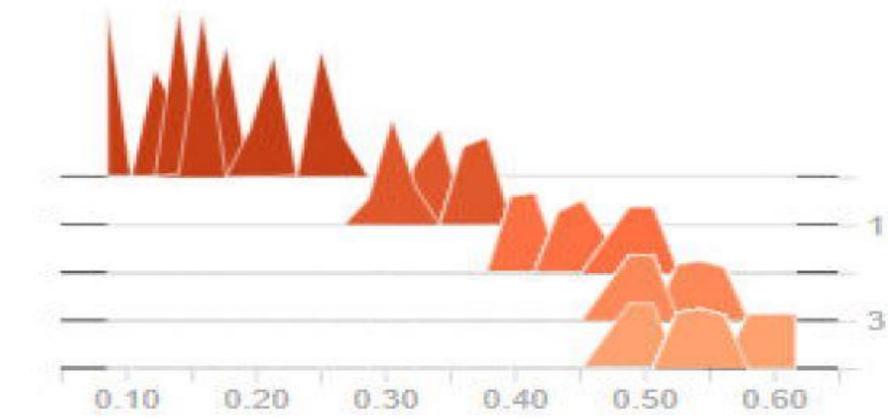


loss

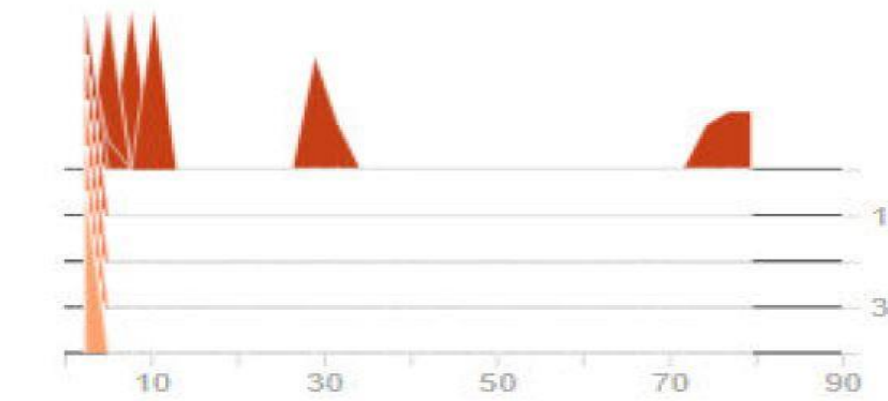


Histograms

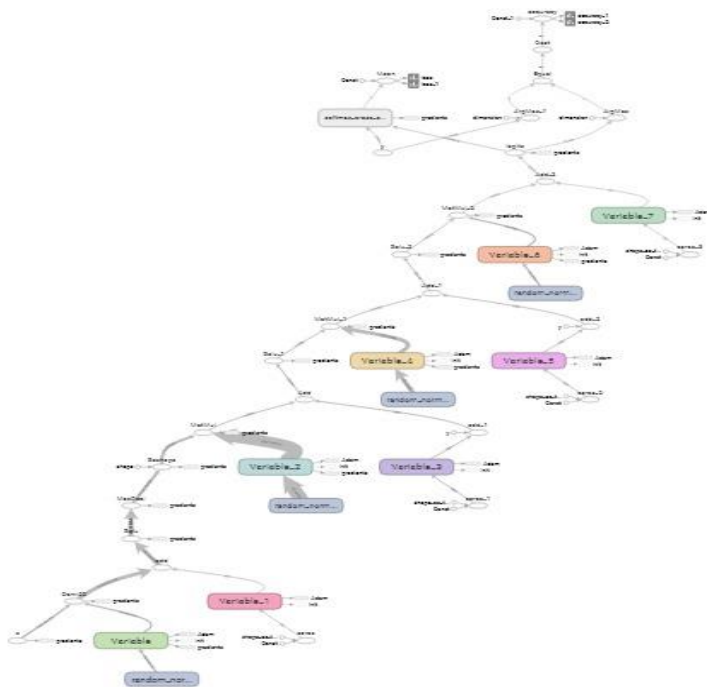
accuracy_1



loss



Graph



Exercise 2: Convolutional Neural Networks with Selu activation function

```
import pickle
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn import preprocessing
#import numpy as np
import random
#import tensorflow as tf
from sklearn.preprocessing import LabelBinarizer

# This function returns the number of labels in the cifar10 dataset
def _load_label_names():

    return ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# This function takes the batches from the cifar10 folder and convert each batch size 3072 to the transformation of depth,length,
def load_cifar10_batch(cifar10_dataset_folder_path, batch_id):

    with open(cifar10_dataset_folder_path + '/data_batch_' + str(batch_id), mode='rb') as file:
        batch = pickle.load(file, encoding='latin1')

    features = batch['data'].reshape((len(batch['data']), 3, 32, 32)).transpose(0, 2, 3, 1)
    labels = batch['labels']

    return features, labels

#This function is used to display the image statistics based on the batch_id and sample_id
#If the batch_id and sample_id is 4 and 1 respectively, it prints the first mage from training batch 4 along with its statistics
def display_stats(cifar10_dataset_folder_path, batch_id, sample_id):
```

#This function is used to display the image statistics based on the batch_id and sample_id
#If the batch_id and sample_id is 4 and 1 respectively, it prints the first mage from training batch 4 along with its statistics
def display_stats(cifar10_dataset_folder_path, batch_id, sample_id):

```

    batch_ids = list(range(1, 6))

    if batch_id not in batch_ids:
        print('Batch Id out of Range. Possible Batch Ids: {}'.format(batch_ids))
        return None

    features, labels = load_cifar10_batch(cifar10_dataset_folder_path, batch_id)

    if not (0 <= sample_id < len(features)):
        print('{} samples in batch {}. {} is out of range.'.format(len(features), batch_id, sample_id))
        return None

    print('\nStats of batch {}'.format(batch_id))
    print('Samples: {}'.format(len(features)))
    print('Label Counts: {}'.format(dict(zip(*np.unique(labels, return_counts=True))))))
    print('First 20 Labels: {}'.format(labels[:20]))

    sample_image = features[sample_id]
    sample_label = labels[sample_id]
    label_names = _load_label_names()

    print('\nExample of Image {}'.format(sample_id))
    print('Image - Min Value: {} Max Value: {}'.format(sample_image.min(), sample_image.max()))
    print('Image - Shape: {}'.format(sample_image.shape))
    print('Label - Label Id: {} Name: {}'.format(sample_label, label_names[sample_label]))
    plt.axis('off')
    plt.imshow(sample_image)

```

#This function is used to do preprocessing stuff like normalizing the features and one hot encoding the labels and the result is

def _preprocess_and_save(normalize, one_hot_encode, features, labels, filename):

```

    features = normalize(features)
    labels = one_hot_encode(labels)

    pickle.dump((features, labels), open(filename, 'wb'))

```

def preprocess_and_save_data(cifar10_dataset_folder_path, normalize, one_hot_encode):

```

    n_batches = 5
    valid_features = []
    valid_labels = []

    for batch_i in range(1, n_batches + 1):
        features, labels = load_cifar10_batch(cifar10_dataset_folder_path, batch_i)
        validation_count = int(len(features) * 0.1)

        # Preprocess and save a batch of training data
        _preprocess_and_save(
            normalize,
            one_hot_encode,
            features[:-validation_count],
            labels[:-validation_count],
            'preprocess_batch_' + str(batch_i) + '.p')

        # Use a portion of training batch for validation
        valid_features.extend(features[-validation_count:])
        valid_labels.extend(labels[-validation_count:])

    # Preprocess and Save all validation data
    _preprocess_and_save(
        normalize,
        one_hot_encode,
        np.array(valid_features),
        np.array(valid_labels),
        'preprocess_validation.p')

    with open(cifar10_dataset_folder_path + '/test_batch', mode='rb') as file:
        batch = pickle.load(file, encoding='latin1')

    # Load the testing data

    test_features = batch['data'].reshape((len(batch['data']), 3, 32, 32)).transpose(0, 2, 3, 1)
    test_labels = batch['labels']

    # Preprocess and Save all testing data
    _preprocess_and_save(
        normalize,
        one_hot_encode,
        np.array(test_features),
        np.array(test_labels),
        'preprocess_test.p')

```

```

def batch_features_labels(features, labels, batch_size):

    for start in range(0, len(features), batch_size):
        end = min(start + batch_size, len(features))
        yield features[start:end], labels[start:end]

def load_preprocess_training_batch(batch_id, batch_size):

    filename = 'preprocess_batch_' + str(batch_id) + '.p'
    features, labels = pickle.load(open(filename, mode='rb'))

    # Return the training data in batches of size <batch_size> or less
    return batch_features_labels(features, labels, batch_size)

#Display the prediction for the random test data based on the model saved after training it.
def display_image_predictions(features, labels, predictions):
    n_classes = 10
    label_names = _load_label_names()
    label_binarizer = LabelBinarizer()
    label_binarizer.fit(range(n_classes))
    label_ids = label_binarizer.inverse_transform(np.array(labels))

    fig, axes = plt.subplots(nrows=4, ncols=2)
    fig.tight_layout()
    fig.suptitle('Softmax Predictions', fontsize=20, y=1.1)

    n_predictions = 3
    margin = 0.05
    ind = np.arange(n_predictions)
    width = (1. - 2. * margin) / n_predictions

    for image_i, (feature, label_id, pred_indices, pred_values) in enumerate(zip(features, label_ids, predictions.indices, predi
        pred_names = [label_names[pred_i] for pred_i in pred_indices]
        correct_name = label_names[label_id]

        axes[image_i][0].imshow(feature)
        axes[image_i][0].set_title(correct_name)
        axes[image_i][0].set_axis_off()

        axes[image_i][1].barh(ind + margin, pred_values[::-1], width)
        axes[image_i][1].set_yticks(ind + margin)
        axes[image_i][1].set_yticklabels(pred_names[::-1])
        axes[image_i][1].set_xticks([0, 0.5, 1.0])

#This is the folder path where i downloaded and extracted my cifar10 dataset
cifar10_dataset_folder_path = 'D:\\Uni Hildesheim\\4th semester\\DDA Lab\\lab 7\\cifar-10-python\\cifar-10-batches-py'

```

```

def normalize(x):

    maximum = np.max(x)
    minimum = np.min(x)
    return (x - minimum) / (maximum - minimum)

def one_hot_encode(x):

    nx = np.max(x) + 1
    return np.eye(nx)[x]

preprocess_and_save_data(cifar10_dataset_folder_path, normalize, one_hot_encode)
valid_features, valid_labels = pickle.load(open('preprocess_validation.p', mode='rb'))

#Defining the feature tensor
def neural_net_image_input(image_shape):

    return tf.placeholder(
        tf.float32,
        [None, image_shape[0], image_shape[1], 3],
        name='x'
    )

#Defining the Label tensor
def neural_net_label_input(n_classes):

    return tf.placeholder(
        tf.float32,
        [None, n_classes],
        name='y'
    )

tf.reset_default_graph()

```



```
#Defining the convolutional neural network with number of output, kernel size, strides, and pool side and padding
def conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize, pool_strides):
```

```
    # Tensorflow has the crappiest API I've ever seen
    input_depth = x_tensor.get_shape().as_list()[-1]
    W = tf.Variable(tf.random_normal(
        [conv_ksize[0], conv_ksize[1], input_depth, conv_num_outputs],
        stddev=0.1
    ))
    b = tf.Variable(tf.zeros(conv_num_outputs))
    conv = tf.nn.conv2d(x_tensor, W, [1, conv_strides[0], conv_strides[1], 1], 'SAME') + b
    conv = tf.nn.relu(conv)
    return tf.nn.max_pool(
        conv,
        [1, pool_ksize[0], pool_ksize[1], 1],
        [1, pool_strides[0], pool_strides[1], 1],
        'SAME')

```

```
def flatten(x_tensor):
```

```
    shape = x_tensor.get_shape().as_list()
    return tf.reshape(x_tensor, [-1, np.prod(shape[1:])])

```

```
#Defined a Fully connected layer with relu activation function
```

```
def fully_conn(x_tensor, num_outputs):
```

```
    shape = x_tensor.get_shape().as_list()
    W = tf.Variable(tf.random_normal([shape[-1], num_outputs], stddev=0.1))
    b = tf.Variable(tf.zeros(num_outputs)) + 0.1
    return tf.nn.relu(tf.add(tf.matmul(x_tensor, W), b))

```

```
#defining the function for output of each network
```

```
def output(x_tensor, num_outputs):
```

```
    shape = x_tensor.get_shape().as_list()
    W = tf.Variable(tf.random_normal([shape[-1], num_outputs]))
    b = tf.Variable(tf.zeros(num_outputs))
    return tf.add(tf.matmul(x_tensor, W), b)

```

```
def conv_net(x):
```

```
    # conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize, pool_strides)
    tmp = conv2d_maxpool(x, 64, [3, 3], [1, 1], [3, 3], [2, 2])

    #Batch normalization
    tf.nn.batch_normalization(tmp, mean=tf.zeros([64]), variance=tf.ones([64]), variance_epsilon=0.01, scale=None, offset=None)
    #Another maxpooling function

    tmp = tf.nn.max_pool(tmp, [1, 3, 3, 1], [1, 2, 2, 1], 'SAME')

    tmp = flatten(tmp)
    #one fully connected layer fc1
    tmp = fully_conn(tmp, 384)

    #second fully connected layer fc2
    tmp = fully_conn(tmp, 192)

    return output(tmp, 10)

```

```
tf.reset_default_graph()
```

```
# Inputs
```

```
x = neural_net_image_input((32, 32, 3))
y = neural_net_label_input(10)

```

```
# Model
```

```
logits = conv_net(x)
```

```
# Name logits Tensor, so that it can be loaded from disk after training
```

```
logits = tf.identity(logits, name='logits')
```

```
# Loss and Optimizer
```

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))
optimizer = tf.train.AdamOptimizer().minimize(cost)

```

```
# Accuracy
```

```
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')
tf.summary.histogram("loss", cost)
tf.summary.scalar("loss", cost)
# Create a summary to monitor accuracy tensor
tf.summary.histogram("accuracy", accuracy)
tf.summary.scalar("accuracy", accuracy)
# Merge all summaries into a single op
merged_summary_op = tf.summary.merge_all()

```

```

def conv_net(x):

    # conv2d_maxpool(x_tensor, conv_num_outputs, conv_ksize, conv_strides, pool_ksize, pool_strides)
    tmp = conv2d_maxpool(x, 64, [3, 3], [1, 1], [3, 3], [2, 2])

    #Batch normalization
    tf.nn.batch_normalization(tmp,mean=tf.zeros([64]),variance=tf.ones([64]),variance_epsilon=0.01,scale=None,offset=None)
    #Another maxpooling function

    tmp = tf.nn.max_pool(tmp,[1, 3, 3, 1],[1, 2, 2, 1],'SAME')

    tmp = flatten(tmp)
    #one fully connected layer fc1
    tmp = fully_conn(tmp, 384)

    #second fully connected layer fc2
    tmp = fully_conn(tmp, 192)

    return output(tmp, 10)

tf.reset_default_graph()

# Inputs
x = neural_net_image_input((32, 32, 3))
y = neural_net_label_input(10)

# Model
logits = conv_net(x)

# Name logits Tensor, so that it can be loaded from disk after training
logits = tf.identity(logits, name='logits')

# Loss and Optimizer
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))
optimizer = tf.train.AdamOptimizer().minimize(cost)

# Accuracy
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32), name='accuracy')
tf.summary.histogram("loss",cost)
tf.summary.scalar("loss", cost)
# Create a summary to monitor accuracy tensor
tf.summary.histogram("accuracy", accuracy)
tf.summary.scalar("accuracy", accuracy)
# Merge all summaries into a single op
merged_summary_op = tf.summary.merge_all()

```

WARNING:tensorflow:From <ipython-input-1-06b5b4e2ae54>:278: softmax_cross_entropy_with_logits (from tensorflow.python.ops.nn_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow into the labels input on backprop by default.

See [@tf.nn.softmax_cross_entropy_with_logits_v2](#).

Training...

```

Epoch 1, CIFAR-10 Batch 1: Cost = 80.84798431396484 - Validation Accuracy = 0.17839999496936798
Epoch 1, CIFAR-10 Batch 2: Cost = 20.6914119720459 - Validation Accuracy = 0.18320000171661377
Epoch 1, CIFAR-10 Batch 3: Cost = 8.855873107910156 - Validation Accuracy = 0.25099998712539673
Epoch 1, CIFAR-10 Batch 4: Cost = 8.763263702392578 - Validation Accuracy = 0.2669999897480011
Epoch 1, CIFAR-10 Batch 5: Cost = 5.199768543243408 - Validation Accuracy = 0.3296000063419342
Epoch 2, CIFAR-10 Batch 1: Cost = 4.5552825927734375 - Validation Accuracy = 0.33559998869895935
Epoch 2, CIFAR-10 Batch 2: Cost = 3.9713339805603027 - Validation Accuracy = 0.35519999265670776
Epoch 2, CIFAR-10 Batch 3: Cost = 3.5299386978149414 - Validation Accuracy = 0.36160001158714294
Epoch 2, CIFAR-10 Batch 4: Cost = 3.0858263969421387 - Validation Accuracy = 0.3792000114917755
Epoch 2, CIFAR-10 Batch 5: Cost = 2.8657493591308594 - Validation Accuracy = 0.38999998569488525
Epoch 3, CIFAR-10 Batch 1: Cost = 2.725909948348999 - Validation Accuracy = 0.38580000400543213
Epoch 3, CIFAR-10 Batch 2: Cost = 2.603868007659912 - Validation Accuracy = 0.39500001072883606
Epoch 3, CIFAR-10 Batch 3: Cost = 2.544926166534424 - Validation Accuracy = 0.39340001344680786
Epoch 3, CIFAR-10 Batch 4: Cost = 2.435222864151001 - Validation Accuracy = 0.4156000018119812
Epoch 3, CIFAR-10 Batch 5: Cost = 2.32495379447937 - Validation Accuracy = 0.41339999437332153
Epoch 4, CIFAR-10 Batch 1: Cost = 2.1152091026306152 - Validation Accuracy = 0.4334000051021576
Epoch 4, CIFAR-10 Batch 2: Cost = 2.1303391456604004 - Validation Accuracy = 0.42820000648498535
Epoch 4, CIFAR-10 Batch 3: Cost = 2.05873703956604 - Validation Accuracy = 0.4431999921798706
Epoch 4, CIFAR-10 Batch 4: Cost = 1.951916217803955 - Validation Accuracy = 0.448199987411499
Epoch 4, CIFAR-10 Batch 5: Cost = 2.0254485607147217 - Validation Accuracy = 0.453000009059906
Epoch 5, CIFAR-10 Batch 1: Cost = 1.832728624343872 - Validation Accuracy = 0.46860000491142273
Epoch 5, CIFAR-10 Batch 2: Cost = 1.959446668624878 - Validation Accuracy = 0.4431999921798706
Epoch 5, CIFAR-10 Batch 3: Cost = 1.8937727212905884 - Validation Accuracy = 0.46480000019073486
Epoch 5, CIFAR-10 Batch 4: Cost = 1.6802111864089966 - Validation Accuracy = 0.47839999198913574
Epoch 5, CIFAR-10 Batch 5: Cost = 1.8028337955474854 - Validation Accuracy = 0.47679999470710754

```


I have run each training batch for 5 epoch, as I don't have GPU and my CPU takes more time for running. The model which I have saved for training will now be used back for predicting the results for any random 4 images from the test dataset. It produces random images with its top-n (here it's 3) probability of that image belonging to those classes and testing accuracy is plotted. My average training accuracy at the end of epoch 5 is 0.47 (i.e. 47%)

```
try:
    if batch_size:
        pass
except NameError:
    batch_size = 64

n_samples = 4
top_n_predictions = 3

def test_model():

    test_features, test_labels = pickle.load(open('preprocess_test.p', mode='rb'))
    loaded_graph = tf.Graph()

    with tf.Session(graph=loaded_graph) as sess:
        # Load model
        loader = tf.train.import_meta_graph(save_model_path + '.meta')
        loader.restore(sess, save_model_path)

        # Get Tensors from loaded model
        loaded_x = loaded_graph.get_tensor_by_name('x:0')
        loaded_y = loaded_graph.get_tensor_by_name('y:0')
        loaded_logits = loaded_graph.get_tensor_by_name('logits:0')
        loaded_acc = loaded_graph.get_tensor_by_name('accuracy:0')

        # Get accuracy in batches for memory limitations
        test_batch_acc_total = 0
        test_batch_count = 0

        for test_feature_batch, test_label_batch in batch_features_labels(test_features, test_labels, batch_size):
            test_batch_acc_total += sess.run(
                loaded_acc,
                feed_dict={loaded_x: test_feature_batch, loaded_y: test_label_batch})
            test_batch_count += 1

        print('Testing Accuracy: {}'.format(test_batch_acc_total/test_batch_count))

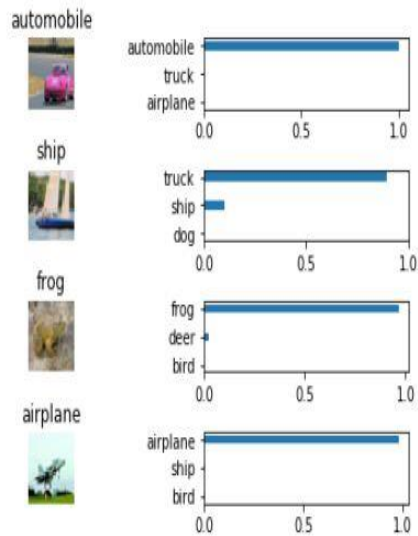
        # Print Random Samples
        random_test_features, random_test_labels = tuple(zip(*random.sample(list(zip(test_features, test_labels)), n_samples)))
        random_test_predictions = sess.run(
            tf.nn.top_k(tf.nn.softmax(loaded_logits), top_n_predictions),
            feed_dict={loaded_x: random_test_features, loaded_y: random_test_labels})
        display_image_predictions(random_test_features, random_test_labels, random_test_predictions)

test_model()
```



```
INFO:tensorflow:Restoring parameters from D:\Uni Hildesheim\4th semester\DDA Lab\lab 7\cifar-10-python\cifar-10-batches-py\image_classification
Testing Accuracy: 0.489899554848671
```

Softmax Predictions




My testing accuracy was 0.48 or (48%) for cifar-10 dataset using selu as activation function without any batch normalization. The above are the prediction of random 4 images and top 3 prediction with respect to each images was predicted correctly or not. For this example image 1, image 3 and image 4 were predicted properly with SoftMax probability > 0.5 , for 2nd image it was misclassified (expected ship but displayed as truck)


Tensorboard Visualization

Scalar

accuracy_2

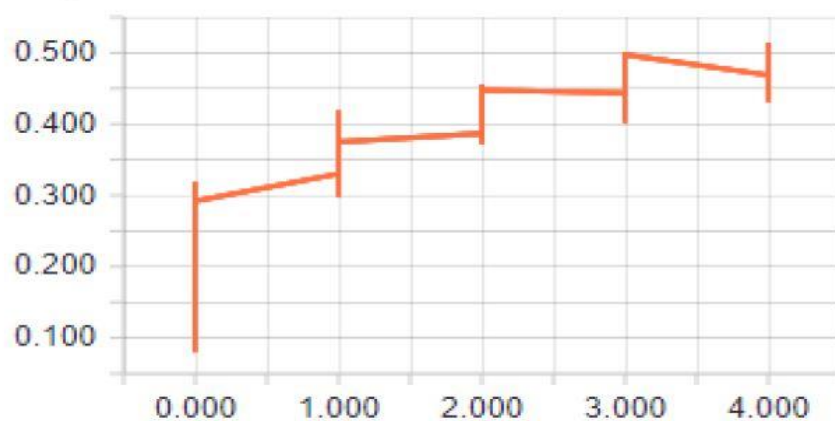


loss_1




Distribution

accuracy_1

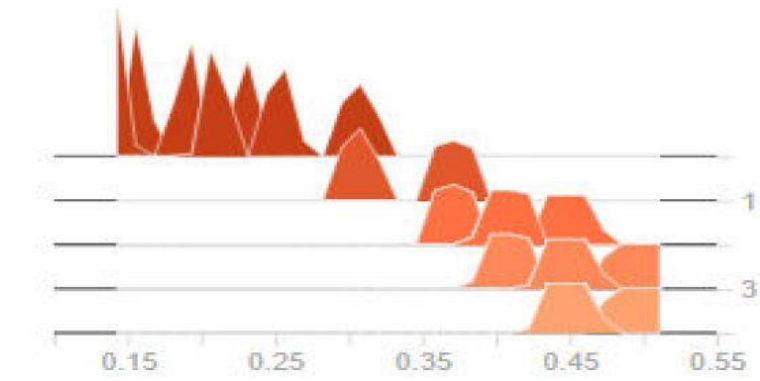


loss

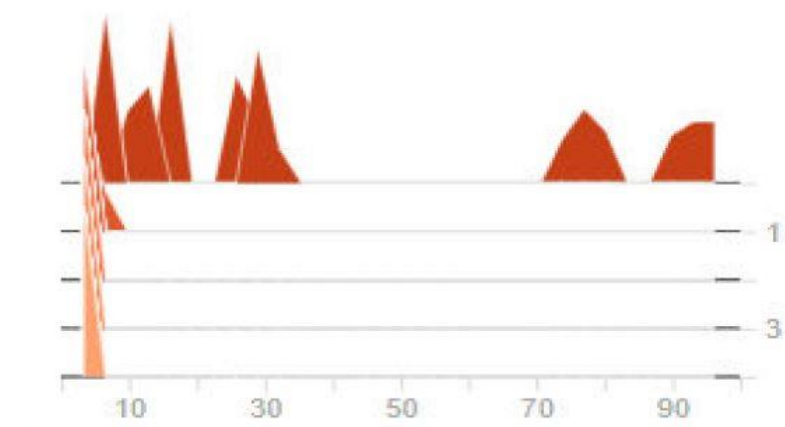


Histogram

accuracy_1



loss



Graph

